



香港中文大學

The Chinese University of Hong Kong

Presented by

Jieming Zhu

jmzhu@cse.cuhk.edu.hk

Scaling Service-oriented Applications into Geo-distributed Clouds

Jieming Zhu, Zibin Zheng, Yangfan Zhou, and Michael R. Lyu

Mar. 25, 2013



Outline

- ▶ Introduction
- ▶ Challenges
- ▶ System Architecture
- ▶ Deployment Approach
- ▶ Experiments
- ▶ Conclusions & Future Work



Introduction

- ▶ **Cloud Computing Systems**
 - Scales up
 - Inter-connection



Amazon Elastic Compute Cloud (Amazon EC2) - Beta



TAP INTO THE
POWER OF NETWORK.COM



MOSSO
the hosting cloud






Introduction

▶ Geo-distributed Clouds

- Many datacenters around the world
- E.g., Amazon: Virginia, Oregon, California, Ireland, Singapore, Tokyo, Sydney, S~ao Paulo, etc.
- Inter-Cloud / Cloud Federation



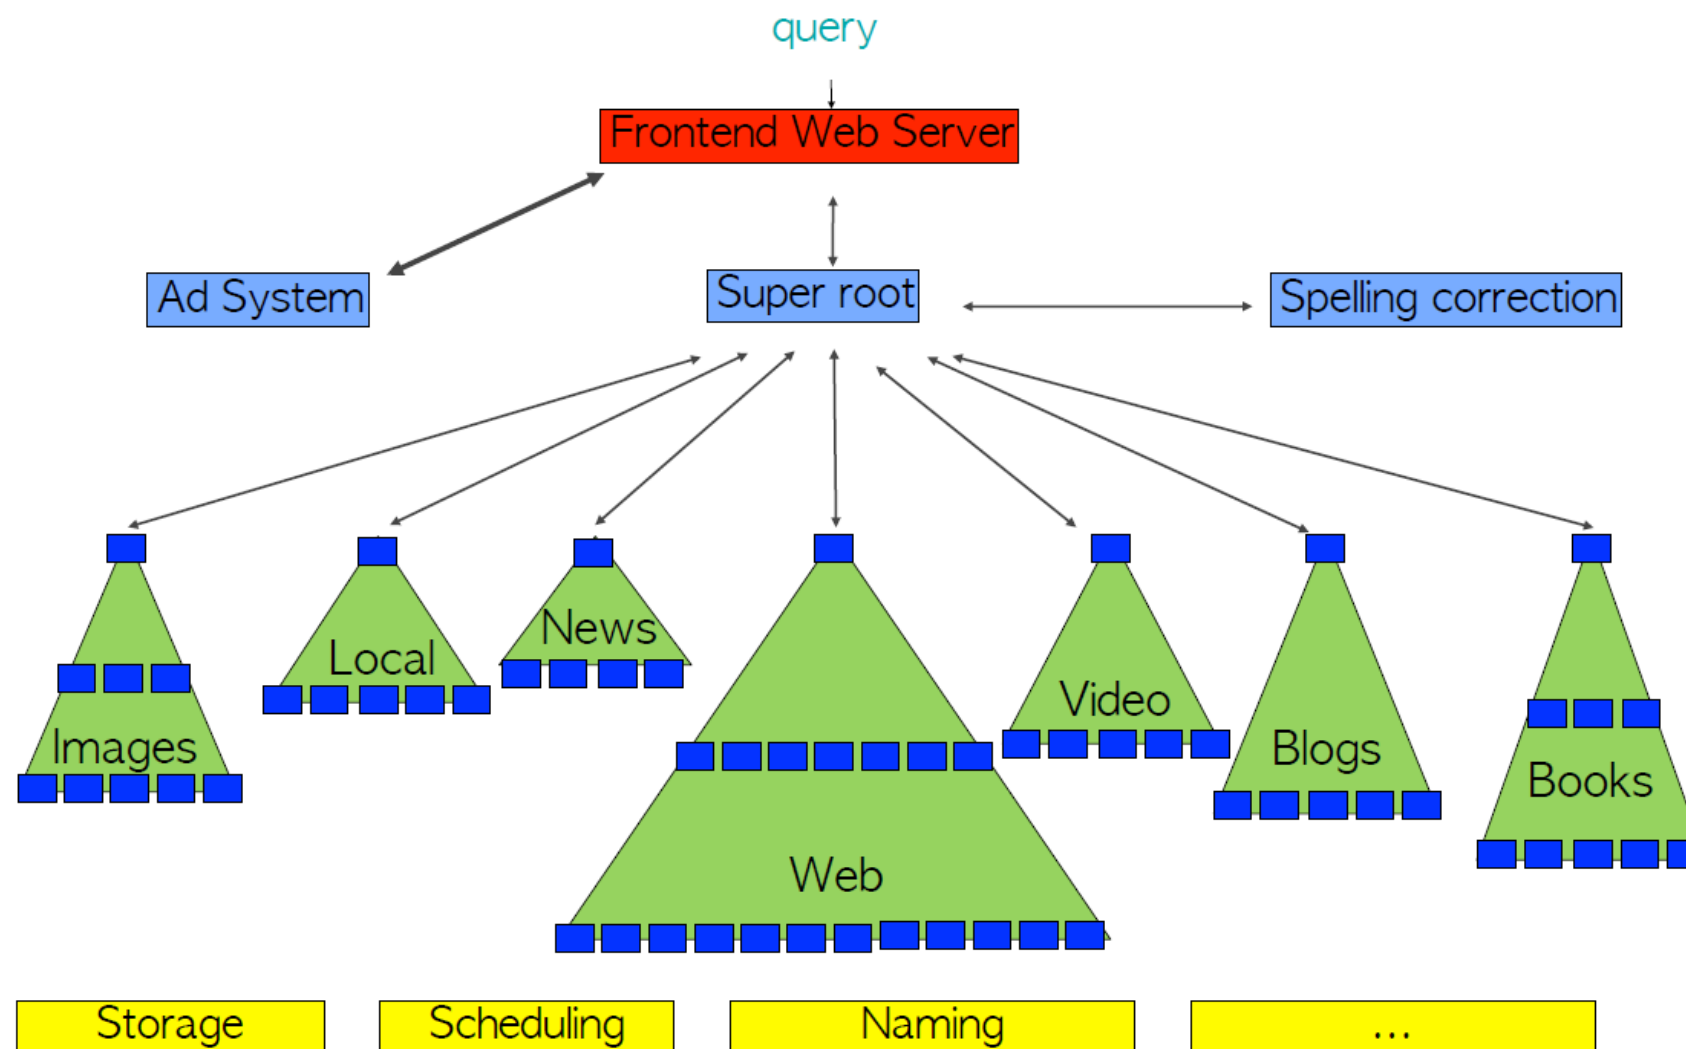
[Figure from Google]



Introduction

► Cloud-based Applications

- Large-scale and complex in business logic
- Service-oriented Architecture
- Service dependency



[Figure from Google]



Challenges

▶ **Opportunities:**

- Take advantage of geo-diversity to
- Optimize performance: E.g., lower latency
- Minimize the operational cost

▶ **Challenges:**

- Dynamic service demand
- Dynamic pricing
- Time-varying communication latencies
- Geographical distribution of end users
- Data sharing and interdependency between services in an application



Application Deployment

- ▶ 1) Which data centers should be selected for each service deployment?
- ▶ 2) How many service instances should be replicated for each service in a data center?



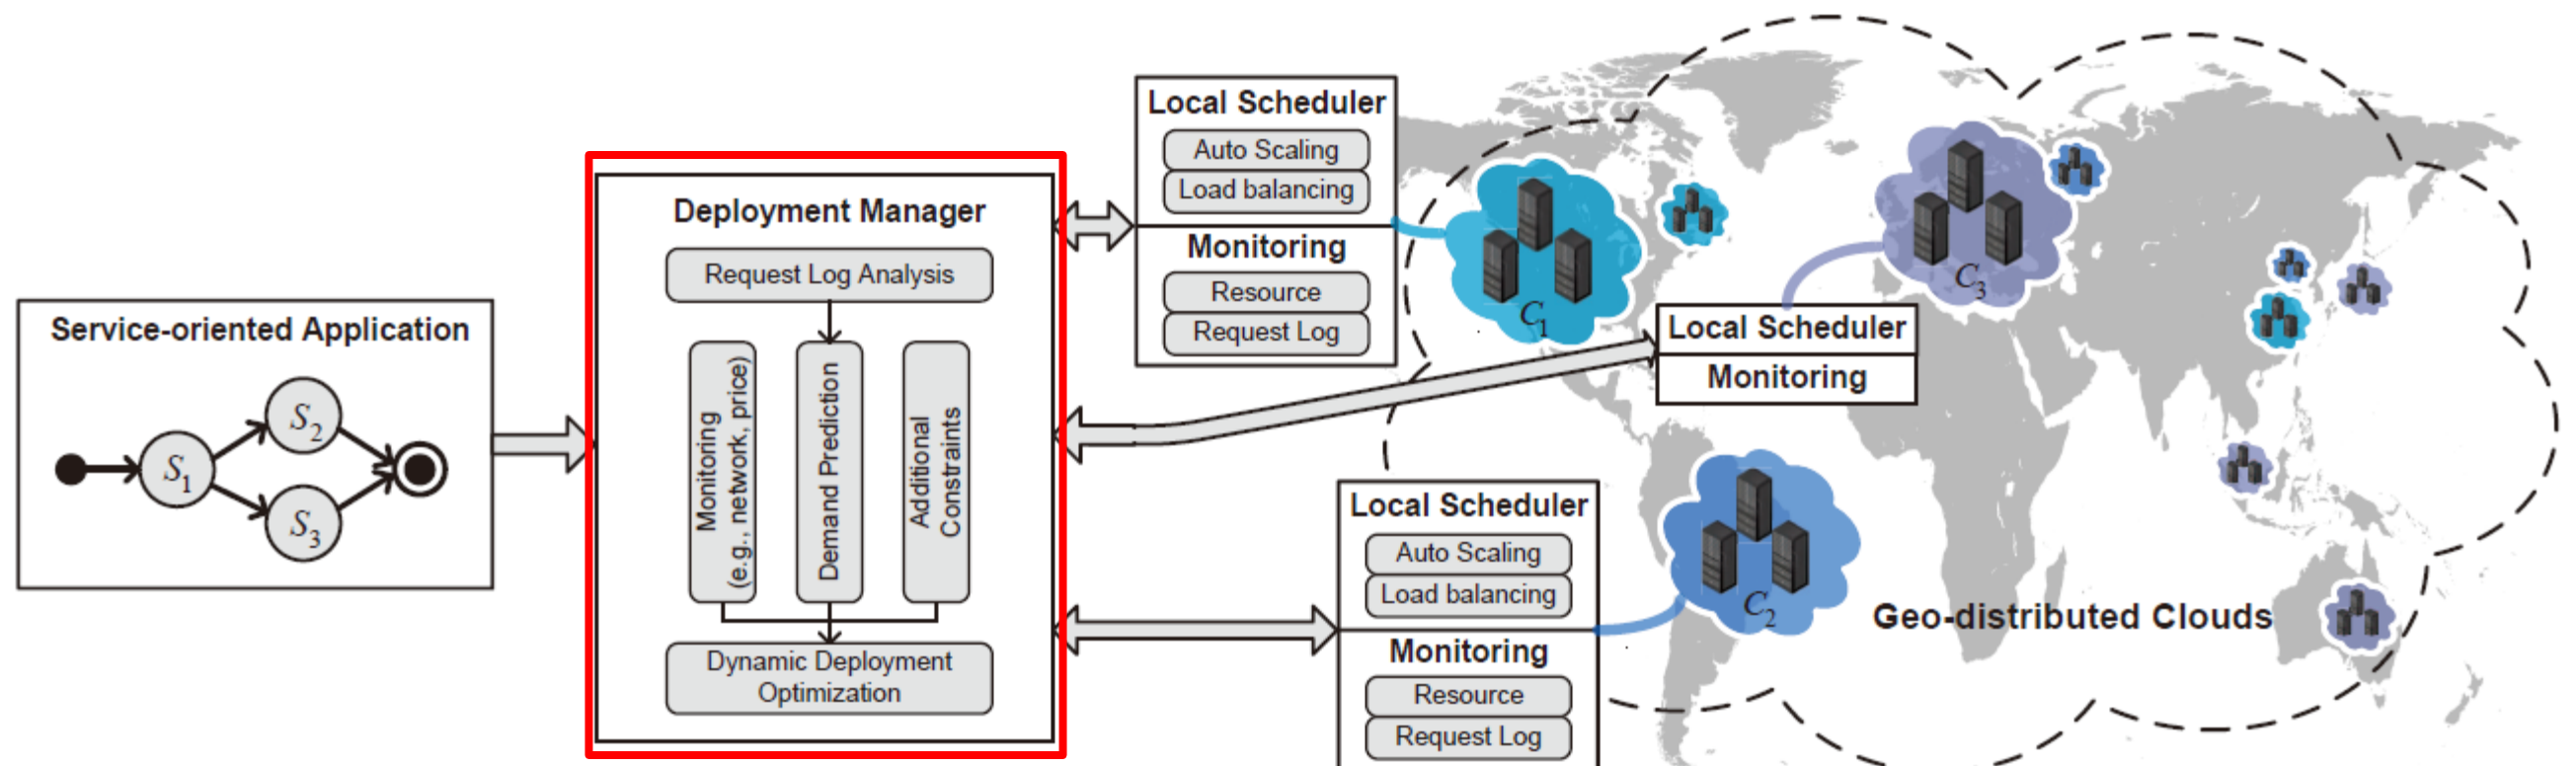
System Architecture



System Architecture

► Deployment Manager

- To determine the locations of each service in the geo-distributed clouds.
- Monitoring: E.g., network, price
- Demand prediction
- Dynamic deployment optimization

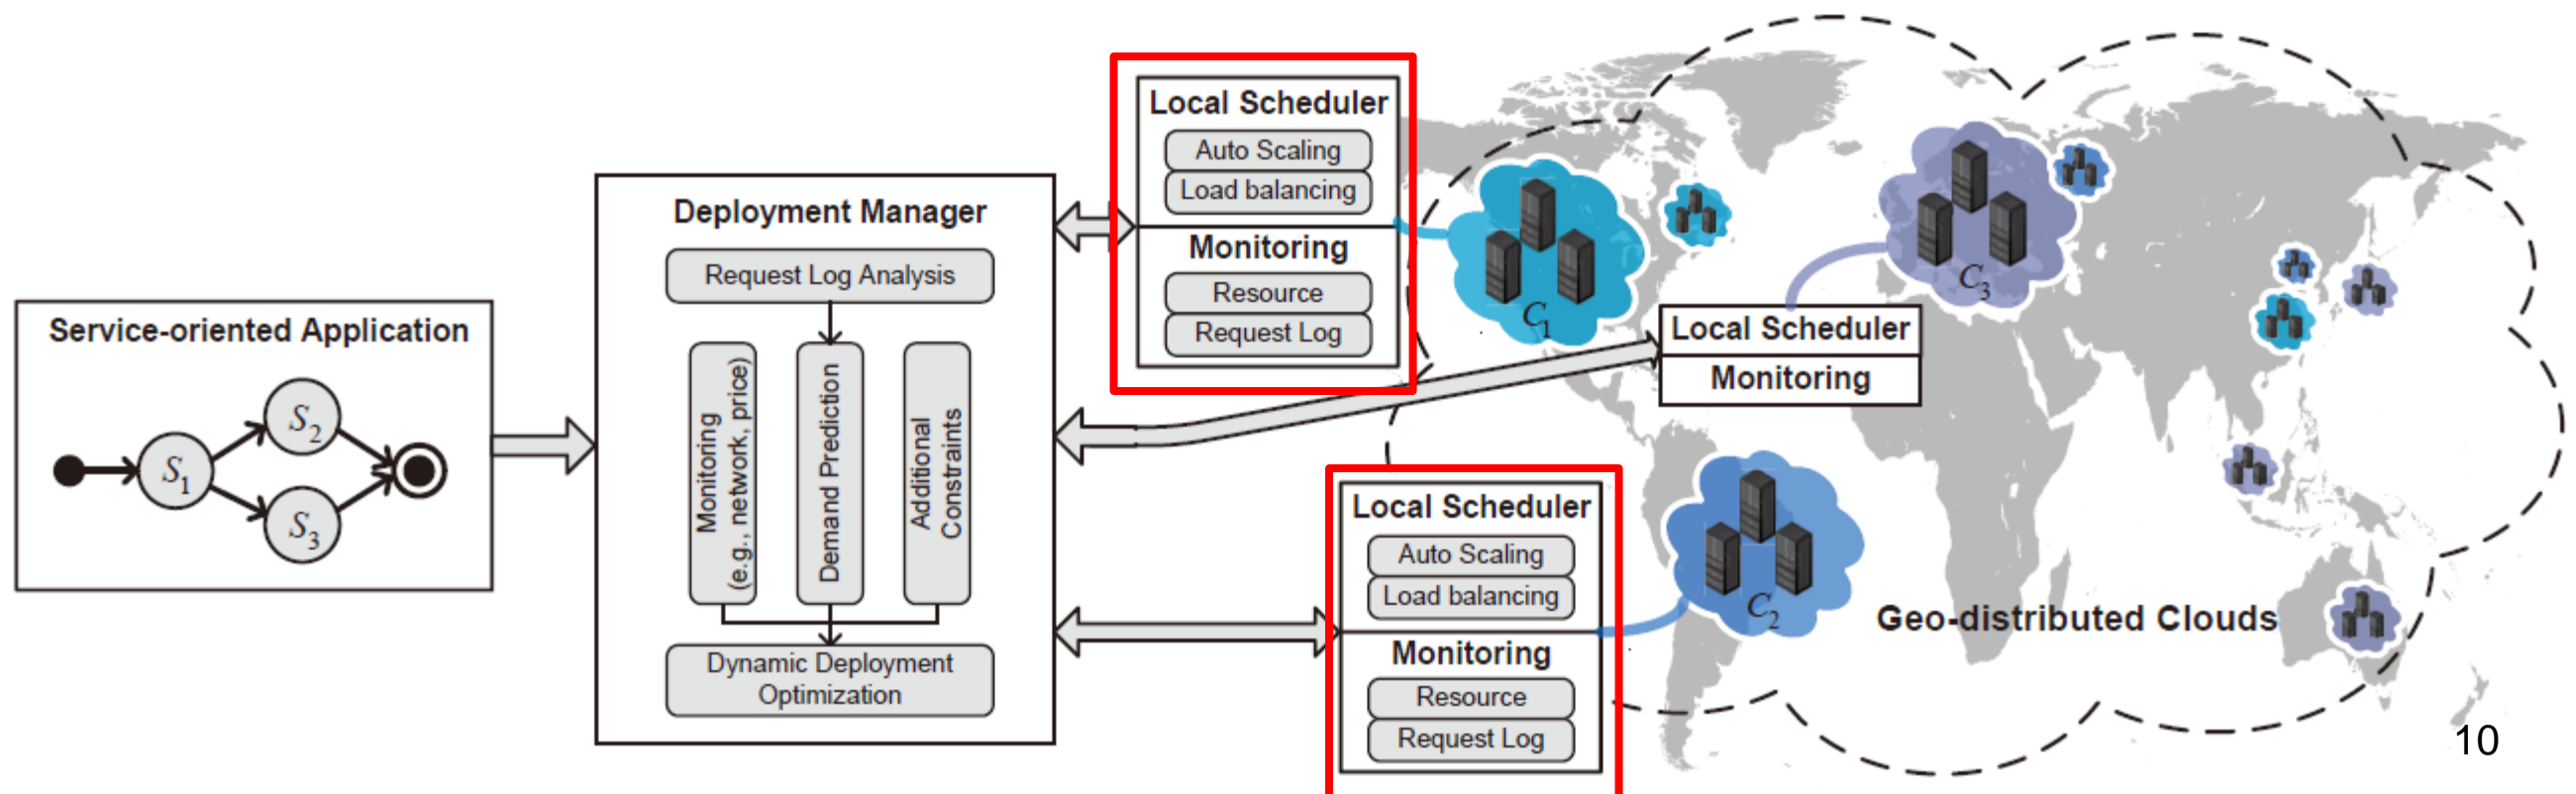




System Architecture

► Local Scheduler

- To scale service instances according to the dynamic request workload and route the service requests with load balancing
- E.g.,
 - **Auto scaling:** <http://aws.amazon.com/autoscaling/>
 - **Elastic load balancing services:** <http://aws.amazon.com/elasticloadbalancing/>





Deployment Approach



Deployment Model

- ▶ **Objective: Minimize the user-perceived average latency**
- ▶ **Constraints:**
 - Operational cost budget \longrightarrow number of service instances K_j
 - Available data center candidates

Model 1 Service Deployment Model across Data Centers

$$\min \sum_{i=1}^N \left(\sum_{j=1}^M f_{ij} \cdot \min_{c_j^m \in C_j} d(i, c_j^m) + \sum_{j=1}^M \sum_{\substack{k=1 \\ k \neq j}}^M f_{jk}^i \cdot \min_{\substack{c_j^m \in C_j \\ c_k^n \in C_k}} d(c_j^m, c_k^n) \right) \quad (1)$$

s.t.

$$C_j \subseteq C, \quad \forall j = 1, 2, \dots, M \quad (2)$$

$$|C_j| = K_j, \quad \forall j = 1, 2, \dots, M \quad (3)$$

TABLE I. NOTATIONS OF DEPLOYMENT MODEL


| Notations | Descriptions |
|---------------|---|
| N | Number of users |
| M | Number of service components in an application |
| f_{ij} | Frequency of invocations between user i and service j |
| c_j^m | The m -th datacenter deployed by service j |
| $d(i, c_j^m)$ | Latency between user i and datacenter c_j^m |
| f_{jk}^i | Frequency of invocations between service j and service k for user i |
| C_j | Deployment strategy of service j (the selected data centers for service j) |
| C | The set of candidate datacenters |
| K_j | The number of instances of service j |

TABLE I. NOTATIONS OF DEPLOYMENT MODEL



Genetic Algorithm Design

▶ Gene Representation



- ▶ Fitness function
- ▶ Parents selection
- ▶ Crossover
- ▶ Mutation

Algorithm 1: Genetic Algorithm based Deployment Algorithm

Input: Fitness function, VMs candidates

Output: The best objective value *Latency*, the deployment solution C_k

```
1 Generate initial population;
2 Compute the fitness value;
3 gen = 0;
4 while gen < MAXGEN do
5     Fitness scaling;
6     Select the parents;
7     Crossover;
8     Mutation;
9     Compute the fitness value of offsprings;
10    Reinsertion;
11    gen++;
12 end
13 return Latency,  $C_k$ 
```



Experiments



Experiments

- ▶ **Dataset description**
 - Planetlab nodes distribution



- Dataset:
 - User-DC matrix: 1881×307
 - Inter-DC matrix: 307×307



Experiments

- ▶ **Dataset description**
 - Latency distribution





Fig. 3. The Approximate Distributions of Inter-DC Latencies and User-DC Latencies (DC: Data Center)



Experiments

► Performance Comparison

- **Random:** The services are deployed randomly in K_j data centers.
- **OneDC:** Deploys all the services in one data center with the highest performance of all the candidates.
- **AIIDC:** Simply deploys each service in every data center.
- **Single Deployment:** Optimizes the deployment strategy independently for each single service



(a) Performance Comparison

User logs: randomly generated


Parameters:

$$N = 1881, M = 10, \text{ and } |C| = 100$$




Experiments

- Impact of parameters



(b) Impact of $|C|$



(c) Impact of K_j

- Find more in our paper



Conclusion

► Contributions:

- We propose a general framework to address the problem of dynamic service-oriented application deployment in geo-distributed clouds
- Optimizing the application performance
- Keeping minimal operational cost.
- Exploits service dependencies to optimize the deployment strategy across a set of candidate data centers
- Extensive experiments are conducted based on a large-scale real-world latency dataset
- Dataset download: <http://www.wsdream.net>

► Future work:

- Explore how to improve our deployment model by incorporating the capacity and cost models
- Load balancing strategies

Thank you!

Q & A